

Global Supply Chain Performance

When the world's shipping lanes narrowed

Data Visualization · Summer 2026 · Final Individual Project

Between late 2023 and 2024, two of the four critical passages in global maritime trade constricted at the same time. Attacks on shipping in the Red Sea pushed carriers away from the Suez Canal and around the Cape of Good Hope, adding roughly ten days to the Asia–Europe voyage. Simultaneously, a drought in Panama forced the canal authority to cut daily transit slots.

Both events were widely reported. What makes them analytically interesting is that we can see them: the IMF's PortWatch platform publishes daily vessel movements derived from satellite AIS signals, so the response of the global fleet is directly observable rather than inferred.

This notebook asks ten analytical questions of that record. The through-line: **the disruption changed the routing of global trade far more than its total volume — and the ports and economies that absorbed the shift were not randomly distributed.**

Dataset

IMF PortWatch (portwatch.imf.org) — daily port-call and trade-volume estimates for **2,065 ports** and daily transit counts for **28 major chokepoints**, built from AIS signals broadcast by ~90,000 vessels. Published by the International Monetary Fund under open terms.

Attribute type	Columns
Temporal	date , year , month , day — daily, 2019 → present
Spatial	port coordinates, country , ISO3 , 28 named chokepoints
Categorical	five cargo classes, systemic-importance classification
Numerical	port calls, import/export volume, transit capacity

An honest caveat, stated up front. Port calls are observed vessel movements. The trade volumes are *model-based estimates* derived from vessel type, capacity and draft — not customs records. They are reliable for measuring relative change over time and should be read cautiously as absolute values. Every conclusion below is framed in terms of change, not level.

0 · Setup

All figures are built with Plotly and routed through `src/theme.py`, which defines one template for the entire project: the CVD-safe Okabe–Ito palette, muted grey for context with a single highlight colour for focus, no gridlines beyond a whisper-light horizontal set, and takeaway-first titles.

```
In [1]: import sys
        from pathlib import Path

        import numpy as np
        import pandas as pd
        import plotly.express as px
        import plotly.graph_objects as go

        sys.path.insert(0, str(Path.cwd().parent))

        from src import data as D
        from src import theme as T  # importing this sets the default Plotly template

        pd.set_option("display.max_columns", 60)
        pd.set_option("display.width", 160)

        print("Palette (CVD-safe, Okabe-Ito):", ", ".join(T.OKABE_ITO))
```

Palette (CVD-safe, Okabe-Ito): orange, sky, green, yellow, blue, vermillion, purple, black

```
In [2]: # Raw data lives in data/raw/ and is gitignored because of its size.
        # If this cell raises FileNotFoundError, run from the repo root:
        #     python scripts/download_data.py

        chokepoints = D.load_chokepoints_daily()
        ports = D.load_ports_daily()
        ports_ref = D.load_ports_reference()

        print(f"chokepoints : {chokepoints.shape[0]:>10,} rows x {chokepoints.shape[1]} cols")
        print(f"ports       : {ports.shape[0]:>10,} rows x {ports.shape[1]} cols")
        print(f"ports_ref    : {ports_ref.shape[0]:>10,} rows x {ports_ref.shape[1]} cols")
        print(f"\ncoverage     : {ports['date'].min():%Y-%m-%d} to {ports['date'].max():%Y-%m-%d}")
```

```
chokepoints :      77,389 rows x 20 cols
ports       : 5,631,255 rows x 29 cols
ports_ref   :      2,065 rows x 25 cols

coverage    : 2019-01-01 to 2026-06-19
```

1 · Preliminary exploration

This section is *not* part of the ten analytical questions — the brief is explicit that value counts and single-variable distributions don't count. It exists to establish what the data actually contains and where it can mislead us, which is a precondition for trusting anything that follows.

```
In [3]: # Structure and completeness
        summary = pd.DataFrame({
            "dtype": ports.dtypes.astype(str),
```

```

"missing": ports.isna().sum(),
"missing_%": (100 * ports.isna().mean()).round(2),
"n_unique": ports.nunique(),
})
summary

```

Out[3]:

	dtype	missing	missing_%	n_unique
date	datetime64[us]	0	0.0	2727
year	int64	0	0.0	8
month	int64	0	0.0	12
day	int64	0	0.0	31
portid	str	0	0.0	2065
portname	str	0	0.0	2045
country	str	0	0.0	183
ISO3	str	0	0.0	180
portcalls_container	int64	0	0.0	67
portcalls_dry_bulk	int64	0	0.0	74
portcalls_general_cargo	int64	0	0.0	46
portcalls_oro	int64	0	0.0	19
portcalls_tanker	int64	0	0.0	115
portcalls_cargo	int64	0	0.0	152
portcalls	int64	0	0.0	190
import_container	int64	0	0.0	128306
import_dry_bulk	int64	0	0.0	161536
import_general_cargo	int64	0	0.0	23501
import_oro	int64	0	0.0	10088
import_tanker	int64	0	0.0	131469
import_cargo	int64	0	0.0	213176
import	int64	0	0.0	248884
export_container	int64	0	0.0	126848
export_dry_bulk	int64	0	0.0	152151
export_general_cargo	int64	0	0.0	23275
export_oro	int64	0	0.0	11685
export_tanker	int64	0	0.0	129311
export_cargo	int64	0	0.0	209343
export	int64	0	0.0	248778

```
In [4]: # Coverage check: does every port report on every day, or is the panel ragged?
# This matters - a port that stops reporting would look like a collapse in traffic
per_day = ports.groupby("date")["portid"].nunique()
print(f"Ports reporting per day: min {per_day.min():,} median {per_day.median():,} max {per_day.max():,}")

gaps = per_day[per_day < per_day.median() * 0.9]
print(f"Days with unusually thin coverage: {len(gaps)}")
if len(gaps):
    print(gaps.head(10))
```

Ports reporting per day: min 2,065 median 2,065 max 2,065

Days with unusually thin coverage: 0

```
In [5]: # Sanity check the cargo decomposition: do the parts sum to the whole?
parts = ports[D.cargo_columns("portcalls")].sum(axis=1)
diff = (parts - ports["portcalls"]).abs()
print(f"Rows where cargo columns != total: {(diff > 0).sum():,} of {len(ports):,}")
print(f"Largest discrepancy: {diff.max():,}")
print("\nNote: 'portcalls_cargo' is a roll-up of the cargo classes, so it is")
print("excluded from the component list to avoid double counting.")
```

Rows where cargo columns != total: 0 of 5,631,255

Largest discrepancy: 0

Note: 'portcalls_cargo' is a roll-up of the cargo classes, so it is excluded from the component list to avoid double counting.

```
In [6]: # What the chokepoint table looks like
print(f"{chokepoints['chokepoint'].nunique():,} chokepoints tracked:\n")
order = chokepoints.groupby("chokepoint")["n_total"].sum().sort_values(ascending=False)
print(order.to_string())
```

28 chokepoints tracked:

chokepoint	
Taiwan Strait	670133
Korea Strait	621813
Malacca Strait	547948
Bohai Strait	486211
Dover Strait	461141
Gibraltar Strait	363112
Bosporus Strait	263470
Strait of Hormuz	237949
Luzon Strait	196137
Cape of Good Hope	170190
Suez Canal	147472
Makassar Strait	143806
Bab el-Mandeb Strait	142672
Oresund Strait	135313
Tsugaru Strait	127010
Yucatan Channel	125916
Mindoro Strait	124576
Lombok Strait	98301
Panama Canal	85472
Sunda Strait	79908
Kerch Strait	75223
Windward Passage	43424
Ombai Strait	30577
Balabac Strait	29587
Mona Passage	29381
Torres Strait	25189
Magellan Strait	10895
Bering Strait	2155

Q1 · Did traffic lost at Suez reappear at the Cape of Good Hope?

Why this question. It is the foundational one. If aggregate volumes simply fell, the story is a demand shock. If volumes held but moved, the story is about *routing* — longer voyages, more ships needed for the same cargo, higher freight costs. Those are very different economic narratives, and the daily transit record can distinguish them.

What I expected. A sharp fall at Suez and Bab el-Mandeb from mid-December 2023, with a compensating rise at the Cape.

Method note. Chokepoints differ by an order of magnitude in absolute traffic, so raw lines would make the smaller passages invisible. Each series is therefore indexed to its own Jan–Jun 2023 average = 100, making the *shapes* comparable. A 7-day centred mean removes the weekly port cycle.

```
In [7]: FOCUS = ["Suez Canal", "Bab el-Mandeb Strait", "Cape of Good Hope", "Panama Cana

# Chokepoint names vary slightly between releases - match loosely, then report
available = chokepoints["chokepoint"].dropna().unique()
matched = {}
for want in FOCUS:
    key = want.split()[0].lower()
    hit = [a for a in available if key in a.lower()]
```

```

    if hit:
        matched[want] = hit[0]
print("Matched chokepoint names:", matched)

sel = chokepoints[chokepoints["chokepoint"].isin(matched.values())].copy()
sel = sel.groupby(["date", "chokepoint"], as_index=False)["n_total"].sum()
sel["n_smooth"] = (sel.sort_values("date").groupby("chokepoint")["n_total"]
                    .transform(lambda s: D.smooth(s, 7)))

sel = D.index_to_baseline(sel, "n_smooth", "chokepoint", "date",
                          ("2023-01-01", "2023-06-30"))
sel = sel[sel["date"] >= "2022-06-01"]
sel.head()

```

Matched chokepoint names: {'Suez Canal': 'Suez Canal', 'Bab el-Mandeb Strait': 'Bab el-Mandeb Strait', 'Cape of Good Hope': 'Cape of Good Hope', 'Panama Canal': 'Panama Canal'}

Out[7]:

	date	chokepoint	n_total	n_smooth	n_smooth_idx
4988	2022-06-01	Bab el-Mandeb Strait	60	63.714286	86.207964
4989	2022-06-01	Cape of Good Hope	34	39.857143	87.460815
4990	2022-06-01	Panama Canal	36	34.571429	107.656008
4991	2022-06-01	Suez Canal	57	61.285714	83.063050
4992	2022-06-02	Bab el-Mandeb Strait	69	64.428571	87.174421

In [8]:

```

fig = go.Figure()

for name in sel["chokepoint"].unique():
    d = sel[sel["chokepoint"] == name].sort_values("date")
    is_suez = "suez" in name.lower()
    fig.add_trace(go.Scatter(
        x=d["date"], y=d["n_smooth_idx"], mode="lines", name=name,
        line=dict(color=T.HIGHLIGHT if is_suez else T.CONTEXT,
                  width=2.8 if is_suez else 1.6),
    ))

# Re-draw the Cape in the secondary highlight - it is the other half of the stor
cape = sel[sel["chokepoint"].str.contains("Cape", case=False, na=False)]
if not cape.empty:
    fig.add_trace(go.Scatter(
        x=cape["date"], y=cape["n_smooth_idx"], mode="lines",
        name=cape["chokepoint"].iloc[0],
        line=dict(color=T.HIGHLIGHT_2, width=2.8),
    ))


fig.add_hline(y=100, line=dict(color=T.CONTEXT, width=1, dash="dot"))
T.event_band(fig, "2023-12-15", sel["date"].max(), "Red Sea crisis")

T.titled(
    fig,
    "Suez traffic collapsed - and reappeared around the Cape of Good Hope",
    "Daily transit calls, 7-day mean, indexed to each chokepoint's Jan-Jun 2023",
    "Source: IMF PortWatch",
)
fig.update_layout(height=520, yaxis_title="Index (2023 H1 = 100)", xaxis_title="")
fig.show()

```

Suez traffic collapsed - and reappeared around the Cape of Good Hope

Daily transit calls, 7-day mean, indexed to each chokepoint's Jan-Jun 2023 average = 100



What I found. Suez and Bab el-Mandeb collapsed together and never fully recovered: indexed to their own late-2023 baseline (=100), Suez fell to an index of **38** at the trough (a 62% drop) and Bab el-Mandeb to **24** (a 76% drop), and both were still running at only **~50-55** — roughly half their pre-crisis level — in the most recent 60 days of the record. The Cape of Good Hope did not just absorb the loss, it overshoot it: its index rose to **204** in the same recent window, more than double its own baseline. The offset is therefore not proportional — Suez/Bab el-Mandeb lost roughly half their traffic while the Cape gained more than 100%, consistent with each transit now requiring a materially longer voyage (more Cape transits are needed to move the same cargo that one Suez transit used to carry).

What this chart cannot tell us. Transit *counts* are not transit *tonnage*. A fall in vessel numbers with larger average ships would overstate the volume loss. Q2 partly addresses this by splitting out cargo classes.

Q2 · Did every cargo type reroute, or only some?

Why this question. Container lines run fixed schedules, carry high-value goods and face reputational and insurance pressure. Tankers operate on different charter structures. If

the two respond differently, the rerouting decision was commercial rather than purely physical — a more interesting finding than a uniform retreat.

Method note. Small multiples with a shared y-axis. Shared scales are a deliberate choice: independent scales would make every facet look equally dramatic and would be quietly dishonest.

```
In [9]: suzez_name = matched.get("Suez Canal")
suez = chokepoints[chokepoints["chokepoint"] == suzez_name].copy()


long = D.to_long_cargo(suez, "n", ["date"])
long = long.groupby(["date", "cargo"], as_index=False)["value"].sum()
long["value"] = (long.sort_values("date").groupby("cargo")["value"]
                 .transform(lambda s: D.smooth(s, 14)))
long = D.index_to_baseline(long, "value", "cargo", "date",
                           ("2023-01-01", "2023-06-30"))
long = long[long["date"] >= "2022-06-01"]

fig = px.line(long, x="date", y="value_idx", facet_col="cargo", facet_col_wrap=3,
              color_discrete_sequence=[T.HIGHLIGHT])
fig.update_traces(line=dict(width=2.2))
fig.for_each_annotation(lambda a: a.update(text=a.text.split("=")[-1],
                                           font=dict(size=13, color=T.INK)))
fig.add_hline(y=100, line=dict(color=T.CONTEXT, width=1, dash="dot"))
fig.update_yaxes(matches="y")

T.titled(
    fig,
    "Vehicle carriers abandoned the Red Sea hardest; general cargo held on longer",
    "Suez Canal transit calls by vessel class, 14-day mean, indexed to Jan-Jun 2022",
    "Source: IMF PortWatch",
)
fig.update_layout(height=560, showlegend=False)
fig.update_xaxes(title="")
fig.show()
```

Vehicle carriers abandoned the Red Sea hardest; general cargo h

Suez Canal transit calls by vessel class, 14-day mean, indexed to Jan-Jun 2023 = 100



What I found. The initial hypothesis — that container traffic would flee first because it is schedule-driven and reputation-sensitive — is not what the data shows. **Ro-Ro (vehicle carrier) traffic fell hardest by far, down 93%** at the trough, followed by container (-62%) and dry bulk (-58%); tanker (-53%) and general cargo (-40%) fell least. A plausible reading: Ro-Ro cargo is high-value, schedule-sensitive *and* the vessels are comparatively vulnerable to the kind of attacks reported in the Red Sea, giving carriers the strongest incentive to reroute entirely rather than accept the risk. General cargo's comparatively small drop suggests some of that trade is on flexible tramp routes less tied to the Suez corridor. This is a case where the honest finding overturned the working hypothesis, which is itself worth stating plainly.

Q3 · Which chokepoints are structurally volatile, and does volatility cluster?

Why this question. Q1 and Q2 examine one crisis. This one asks a structural question across the whole seven-year record: is instability a property of certain passages, or does it strike arbitrarily? The answer bears on how a shipper should think about route risk in general.

Method note. Coefficient of variation (σ/μ) rather than raw standard deviation, so that busy and quiet passages are comparable.

```
In [10]: stats = (chokepoints.groupby("chokepoint")["n_total"]
            .agg(mean="mean", sd="std", total="sum").reset_index())
stats = stats[stats["mean"] > 1]
stats["cv"] = stats["sd"] / stats["mean"]
stats = stats.sort_values("cv").tail(20)


colors = T.emphasise(len(stats), list(range(len(stats) - 3, len(stats))))

fig = go.Figure()
fig.add_trace(go.Scatter(
    x=stats["cv"], y=stats["chokepoint"], mode="markers",
    marker=dict(size=13, color=colors),
    hovertemplate="<b>{y}</b><br>CV {x:.2f}<extra></extra>",
))
for _, r, c in zip(range(len(stats)), stats.itertuples(), colors):
    fig.add_shape(type="line", x0=0, x1=r.cv, y0=r.chokepoint, y1=r.chokepoint,
        line=dict(color=c, width=2))

T.titled(
    fig,
    "The most volatile passages are a mix of geopolitical flashpoints and thin,
    "Coefficient of variation of daily transit calls, full record - 20 most vari
    "Source: IMF PortWatch",
)
fig.update_layout(height=640, xaxis_title="Coefficient of variation", yaxis_titl
fig.show()
```

The most volatile passages are a mix of geopolitical flashpoints

Coefficient of variation of daily transit calls, full record - 20 most variable chokepoints



What I found. The three highlighted, most-volatile chokepoints are **Kerch Strait, Magellan Strait, and the Cape of Good Hope** (after the lowest-traffic route, the Bering Strait, is excluded as pure statistical noise). Two of the three have an obvious common thread: Kerch sits in the Russia-Ukraine Black Sea war zone and the Cape's volatility is exactly the Red Sea reroute surge measured in Q1 — both are genuine geopolitical shocks. Magellan Strait breaks the pattern, though: it is a low-traffic, seasonally ice-affected passage (only a few vessels a day on average), so a handful of extra transits swings its ratio wildly — a useful counter-example showing that coefficient of variation, taken alone, cannot fully separate "deliberate political disruption" from "small numbers exaggerating an ordinary swing." Notably, the Strait of Hormuz — the passage with the most persistent real-world geopolitical tension — does *not* rank in the top three here, because its baseline traffic is large enough that even its swings during periods of tension are proportionally modest. Volatility clusters partly around geopolitics, but this metric rewards routes that are both risky *and* thin, which is a distinct combination from "important and risky."

Q4 · Did Panama and the Red Sea compound each other?

Why this question. This is the crux. Asia→US-East-Coast cargo has two main options: west through Panama, or east through Suez. Both narrowed within months of each other. If the constraints overlapped in time, shippers had no good alternative — a genuinely unusual situation in modern maritime trade.

```
In [11]: pair = [matched.get("Suez Canal"), matched.get("Panama Canal")]
pair = [p for p in pair if p]

d = chokepoints[chokepoints["chokepoint"].isin(pair)].copy()
d = d.groupby(["date", "chokepoint"], as_index=False)["n_total"].sum()
d["smooth"] = (d.sort_values("date").groupby("chokepoint")["n_total"]
               .transform(lambda s: D.smooth(s, 14)))
d = D.index_to_baseline(d, "smooth", "chokepoint", "date",
                       ("2022-01-01", "2022-12-31"))
d = d[d["date"] >= "2022-01-01"]


fig = go.Figure()
for i, name in enumerate(pair):
    sub = d[d["chokepoint"] == name].sort_values("date")
    fig.add_trace(go.Scatter(
        x=sub["date"], y=sub["smooth_idx"], mode="lines", name=name,
        line=dict(color=[T.HIGHLIGHT, T.HIGHLIGHT_2][i % 2], width=2.6),
    ))

fig.add_hline(y=100, line=dict(color=T.CONTEXT, width=1, dash="dot"))
T.event_band(fig, "2023-07-01", "2024-06-30", "Panama drought", y=1.0)
T.event_band(fig, "2023-12-15", "2024-12-31", "Red Sea crisis", y=0.92)

T.titled(
    fig,
    "Both Asia-to-US-East-Coast routes narrowed at once, leaving no easy alterna
    "Transit calls, 14-day mean, indexed to each chokepoint's 2022 average = 100
    "Source: IMF PortWatch",
)
fig.update_layout(height=520, yaxis_title="Index (2022 = 100)", xaxis_title="")
fig.show()
```

Both Asia-to-US-East-Coast routes narrowed at once, leaving no

Transit calls, 14-day mean, indexed to each chokepoint's 2022 average = 100



What I found. The two windows overlapped for **6.5 months**, from 15 December 2023 to 30 June 2024. For that entire stretch, Panama Canal transits ran at their drought-restricted trough — **32.1 transits/day before the drought down to 15.0/day at the low point, a 53% fall** — at the same time Suez/Bab el-Mandeb traffic was down 62-76% (Q1). Asia-to-US-East-Coast cargo lost its two shortest options simultaneously for over half a year: the Panama shortcut and the Suez-plus-Cape-alternative pairing. That is the crux of the story — this was not one canal absorbing a shock while the other stood by, it was the two principal alternatives to a Cape diversion both constrained at once.

Q5 · Which ports won and lost, and where are they?

Why this question. Aggregate rerouting has to land somewhere physical. Naming the ports turns an abstract shipping story into a concrete one, and the geography is the test of whether the pattern makes sense.

Method note. Mean weekly calls in Jan–Sep 2024 versus Jan–Sep 2023, on a like-for-like seasonal window. Ports below a minimum baseline are excluded so that a port going from 1 call to 6 doesn't top the chart at +500%.

```
In [12]: p = ports.copy()
p["week"] = p["date"].dt.to_period("W").dt.start_time
weekly = p.groupby(["week", "portid", "portname", "country", "ISO3"],
                    as_index=False)["portcalls_container"].sum()

changes = D.pct_change_between(
    weekly, "portcalls_container", "portname", "week",
    before=("2023-01-01", "2023-09-30"),
    after=("2024-01-01", "2024-09-30"),
    min_base=10.0,
)
print(f"{len(changes):,} ports with a meaningful container baseline")
changes.head()
```

207 ports with a meaningful container baseline

```
Out[12]:
```

	portname	before	after	abs_change	pct_change
0	Mersin	21.282051	29.025	7.742949	36.382530
1	Alexandria	19.384615	24.800	5.415385	27.936508
2	Ulsan	16.923077	20.525	3.601923	21.284091
3	Mumbai-Jawaharlal Nehru (Nhava Sheva)	44.974359	54.475	9.500641	21.124572
4	Taipei	25.205128	30.325	5.119872	20.312818


```
In [13]: top = pd.concat([changes.head(15), changes.tail(15)]).sort_values("pct_change")

fig = go.Figure(go.Bar(
    x=top["pct_change"], y=top["portname"], orientation="h",
    marker=dict(color=top["pct_change"],
                colorscale=[[0.0, T.HIGHLIGHT_2], [0.5, "#F2F2F2"], [1.0, T.HIGH
                cmid=0, line=dict(width=0)),
    hovertemplate="<b>%{y}</b><br>{x:+.1f}%<extra></extra>",
))
fig.add_vline(x=0, line=dict(color=T.INK_SOFT, width=1))

T.titled(
    fig,
    "Rerouting redrew the map of container traffic port by port",
    "Change in mean weekly container calls, Jan-Sep 2024 vs Jan-Sep 2023",
    "Source: IMF PortWatch",
)
fig.update_layout(height=820, xaxis_title="Change (%)", yaxis_title="")
fig.show()
```

Rerouting redraw the map of container traffic port by port

Change in mean weekly container calls, Jan-Sep 2024 vs Jan-Sep 2023



```
In [14]: # The same finding, mapped - geography is the test of whether the pattern is coh
geo = changes.merge(
    ports[["portname", "portid"]].drop_duplicates("portname"),
    on="portname", how="left",
).merge(
    ports_ref[[c for c in ["portid", "latitude", "longitude"] if c in ports_ref.
        .drop_duplicates("portid"),
        on="portid", how="left",
    ].dropna(subset=["latitude", "longitude"]])
```

```

geo = geo[geo["before"] >= 25]


fig = px.scatter_geo(
    geo, lat="latitude", lon="longitude",
    color="pct_change", size=geo["before"],
    hover_name="portname",
    color_continuous_scale=[[0.0, T.HIGHLIGHT_2], [0.5, "#EDED"], [1.0, T.HIGH
    color_continuous_midpoint=0, size_max=24, projection="natural earth",
)
fig.update_traces(marker=dict(line=dict(width=0), opacity=0.85))

T.titled(
    fig,
    "The winners and losers cluster along the re-routed corridors",
    "Change in weekly container calls, 2024 vs 2023 - bubble size = baseline vol
    "Source: IMF PortWatch",
)
fig.update_layout(height=560,
                    coloraxis_colorbar=dict(title="Change %", thickness=12))
fig.show()

```

The winners and losers cluster along the re-routed corridors

Change in weekly container calls, 2024 vs 2023 - bubble size = baseline volume



What I found. The losers are geographically coherent and point straight at the Red Sea: **King Abdullah Port, Saudi Arabia (-81%), Aqaba, Jordan (-71%), and Jeddah, Saudi Arabia (-37%)** are all Red Sea-facing container gateways that lost the feeder traffic which used to transit the corridor they sit on. The winners are more mixed but still tell a routing

story: **Iskenderun and Mersin, Türkiye (+76% and +36%)** and **El Sokhna, Egypt (+42%)** sit on the Mediterranean side and plausibly picked up cargo re-timed around the disrupted corridor, while **Saint Petersburg, Russia (+69%)** and **Wilhelmshaven, Germany (+48%)** reflect Northern European and Baltic trade patterns largely unrelated to Suez. So the losers are a clean, single-cause story; the winners are a mix of Red-Sea-adjacent gains and unrelated trends, which is worth saying plainly rather than forcing every gainer into the same narrative.

Q6 · Is global container traffic becoming more concentrated?

Why this question. Concentration is a fragility measure. If a growing share of traffic flows through fewer ports, the network has more single points of failure — which is exactly what makes a Suez or Panama event systemic rather than local.

```
In [15]: w = weekly.copy()
w["year"] = w["week"].dt.year

rows = []
for yr, grp in w.groupby("year"):
    shares = grp.groupby("portid")["portcalls_container"].sum()
    shares = shares[shares > 0]
    if len(shares) > 50:
        rows.append({
            "year": yr,
            "hhi": D.herfindahl(shares),
            "top20_share": 100 * shares.nlargest(20).sum() / shares.sum(),
            "n_ports": len(shares),
        })
conc = pd.DataFrame(rows)
conc
```

```
Out[15]:
```

	year	hhi	top20_share	n_ports
0	2018	0.009024	33.357558	719
1	2019	0.007803	30.838549	1353
2	2020	0.007459	30.076675	1321
3	2021	0.007244	29.429530	1279
4	2022	0.007638	30.727204	1276
5	2023	0.007683	30.734519	1264
6	2024	0.007526	30.440135	1269
7	2025	0.007737	31.064465	1254
8	2026	0.007791	31.298883	1184

```
In [16]: fig = go.Figure(go.Scatter(
    x=conc["year"], y=conc["top20_share"], mode="lines+markers",
```

```

    line=dict(color=T.HIGHLIGHT, width=3),
    marker=dict(size=10, color=T.HIGHLIGHT),
))


if len(conc) > 1:
    last = conc.iloc[-1]
    T.annotate(fig, last["year"], last["top20_share"],
               f"{last['top20_share']:.1f}% in {int(last['year'])}", ax=-70, ay=

T.titled(
    fig,
    "The 20 busiest ports handle a persistently large share of global container
    Share of all container port calls handled by the top 20 ports, by year",
    "Source: IMF PortWatch",
)
fig.update_layout(height=440, yaxis_title="Top-20 share (%)", xaxis_title="")
fig.show()

```

The 20 busiest ports handle a persistently large share of global container

Share of all container port calls handled by the top 20 ports, by year



What I found. The top-20 ports' share of global container calls is essentially **flat across the whole record** — 30.9% in 2019 versus 31.3% in 2026, drifting inside a narrow 29–31% band throughout, with no sustained rise or fall even through the 2023–24 crisis. Concentration did not increase because of the disruption, and it has not fallen either — the honest reading is a null result on the trend question, not the rising-fragility story the working title implied. That said, a stable ~30% share sitting in just twenty of the roughly 2,065 ports tracked is itself the finding worth keeping: whatever the trend, twenty hubs consistently carry close to a third of global container-port activity, so the network's exposure to a shock at any one of them does not go away just because the share isn't rising.

Q7 · Does logistics quality buy resilience?

Why this question. The payoff question of the whole project. The World Bank publishes a Logistics Performance Index scoring countries on customs, infrastructure, timeliness and tracking. If high-LPI countries recover port traffic faster after a shock, the index measures something real about resilience — not just steady-state efficiency.

Before running this cell: download the LPI dataset from lpi.worldbank.org and save it as `data/raw/lpi.csv` with at least a country/ISO3 column and an overall score column.

Honest caveat. This is a cross-sectional correlation across ~100 countries with one shock. It is suggestive, not causal, and the notebook should say so.

```
In [17]: LPI_PATH = D.RAW / "lpi.csv"

if not LPI_PATH.exists():
    print("lpi.csv not found - download it from https://lpi.worldbank.org/")
    print("and save to", LPI_PATH)
else:
    lpi = pd.read_csv(LPI_PATH)
    print("LPI columns:", list(lpi.columns))
    # Adjust these two names to match the file you downloaded
    ISO_COL, SCORE_COL = "ISO3", "LPI Score"

    # Recovery half-life per country from the daily port-call series
    cty = (ports.groupby(["date", "ISO3"], as_index=False)["portcalls_container"]
           results = []
    for iso, grp in cty.groupby("ISO3"):
        s = grp.set_index("date")["portcalls_container"].sort_index()
        s = D.smooth(s, 14)
        if s.loc["2023-01-01":"2023-10-01"].mean() < 20:
            continue
        hl = D.recovery_half_life(s, "2023-12-15", ("2023-06-01", "2023-11-30"))
        results.append({"ISO3": iso, "half_life_days": hl,
                        "baseline": s.loc["2023-06-01":"2023-11-30"].mean()})
    rec = pd.DataFrame(results)
    merged = rec.merge(lpi[[ISO_COL, SCORE_COL]], left_on="ISO3",
                       right_on=ISO_COL, how="inner").dropna()
    print(f"\n{len(merged)} countries matched")
    merged.head()
```

LPI columns: ['ISO3', 'Country', 'LPI Score', 'Year']
17 countries matched

```
In [18]: if LPI_PATH.exists() and "merged" in dir() and len(merged):
    fig = px.scatter(
        merged, x=SCORE_COL, y="half_life_days", size="baseline",
        hover_name="ISO3", trendline="ols",
        color_discrete_sequence=[T.HIGHLIGHT_2], size_max=34,
    )
    for tr in fig.data:
        if tr.mode == "lines":
            tr.line.color = T.HIGHLIGHT
            tr.line.width = 2

    T.titled(
```


```

fig,
    "Logistics quality only weakly predicts recovery speed - Singapore is th
    "World Bank LPI score vs days to recover half the post-shock gap in cont
    "Sources: IMF PortWatch; World Bank Logistics Performance Index",
)
fig.update_layout(height=520, xaxis_title="LPI score",
                  yaxis_title="Recovery half-life (days)")
fig.show()
else:
    print("Skipped - add data/raw/lpi.csv to run this question.")

```

Logistics quality only weakly predicts recovery speed - Singapore

World Bank LPI score vs days to recover half the post-shock gap in container calls



What I found. Across the **17 countries** with both a usable container baseline and a matched World Bank LPI score, LPI and recovery half-life are positively correlated ($r = 0.39$) — weak-to-moderate, and in the direction the hypothesis predicted: lower-LPI countries tend to sit at the fast end. But the relationship is not clean. **Singapore has the highest LPI score in the sample (4.3) and the single longest half-life (63 days)** — a direct outlier against the hypothesis, plausibly because Singapore's baseline traffic is itself a Suez/Bab el-Mandeb transshipment hub, so its "recovery" is tangled up with the very disruption being measured rather than a clean test of logistics quality. Meanwhile fast recoveries appear at both ends of the LPI scale (Japan and the USA at the high end, Indonesia at the low end, both around 4-8 days). With $n=17$ and one shared shock, this reads as suggestive at best, not a demonstrated causal link between logistics quality and resilience.

Q8 · Which economies are most exposed to a single chokepoint?

Why this question. Exposure is structural — it exists whether or not a crisis is under way. Identifying it turns the analysis from retrospective into something a policymaker could act on.

Method note. A simple, transparent proxy: for each country, the share of its maritime traffic that plausibly routes through its single most-used chokepoint. Document your assignment rule explicitly — this is a modelling choice, not a measurement, and the marker will want to see you acknowledge that.


```
In [19]: # Country-level maritime activity as the exposure denominator
cty = ports.groupby("ISO3", as_index=False).agg(
    calls=("portcalls", "sum"),
    imports=("import", "sum"),
    exports=("export", "sum"),
)
cty["trade"] = cty["imports"] + cty["exports"]
cty = cty[cty["calls"] > 500]

# Simple, documented proxy for chokepoint dependence: import share of total trade
# Replace with a geographic routing assignment if you have time - and say which
# you used in the notebook text either way.
cty["import_dependence"] = 100 * cty["imports"] / cty["trade"]

fig = px.choropleth(
    cty, locations="ISO3", color="import_dependence",
    color_continuous_scale=T.SEQUENTIAL, projection="natural earth",
)
T.titled(
    fig,
    "Import-dependent economies carry the most chokepoint risk",
    "Seaborne imports as a share of total seaborne trade, full record",
    "Source: IMF PortWatch",
)
fig.update_layout(height=540,
                    coloraxis_colorbar=dict(title="Import share %", thickness=12))
fig.show()
```

Import-dependent economies carry the most chokepoint risk

Seaborne imports as a share of total seaborne trade, full record



What I found. The proxy cleanly separates two very different kinds of exposure. The most import-dependent economies — **French Polynesia (97%), Cayman Islands (96%), Haiti (96%), Puerto Rico (96%), Kenya (95%), Bangladesh (95%)** — are small island territories or coastal economies with limited domestic production and few alternative supply routes, so a single-corridor disruption hits them with no easy substitute. At the other end, the least import-dependent — **Kazakhstan (5%), Russia (8%), Australia (9%), Gabon (9%), Iraq (9%)** — are commodity and resource exporters (oil, gas, minerals) whose seaborne trade is dominated by outbound cargo, so an import-side chokepoint shock would barely register in this measure. The clear limitation: this is a trade-*balance* proxy, not a routing measurement — an import-dependent economy whose imports arrive via several redundant corridors is not actually as exposed as this chart implies, and a proper answer would need port-level chokepoint attribution rather than country-level import share.

Q9 · Do global hubs and domestic ports keep different hours?

Why this question. A genuinely multi-dimensional question: activity by day-of-week, conditioned on the port's systemic classification. If global hubs run closer to round-the-

clock while domestic ports keep a weekly rhythm, that operational difference is visible from orbit — and it affects how quickly each type can absorb a surge.

```
In [20]: p2 = D.add_calendar(ports)

# The Live PortWatch reference table has no ready-made systemic/importance
# label, so build a transparent proxy instead: rank ports by total vessel
# traffic (vessel_count_total) and split at the top 20 - a defensible,
# documented modelling choice, not a claim about IMF methodology.
class_col = "port_class"
if "vessel_count_total" in ports_ref.columns:
    ranked = ports_ref[["portid", "vessel_count_total"]].dropna()
    top_ids = set(ranked.nlargest(20, "vessel_count_total")["portid"])
    ports_ref[class_col] = ports_ref["portid"].apply(
        lambda pid: "Globally systemic (top 20 by vessel traffic)" if pid in top
        else "Regional / domestic")
else:
    class_col = None
print("Classification column:", class_col)

if class_col:
    p2 = p2.merge(ports_ref[["portid", class_col]].drop_duplicates("portid"),
                  on="portid", how="left")
    grp = (p2.groupby([class_col, "dow_name"], as_index=False)["portcalls"].mean
           grp["idx"] = 100 * grp["portcalls"] / grp.groupby(class_col)["portcalls"].tr


    order = ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday",
             "Saturday", "Sunday"]
    pivot = grp.pivot(index=class_col, columns="dow_name", values="idx")[order]

    fig = go.Figure(go.Heatmap(
        z=pivot.values, x=pivot.columns, y=pivot.index,
        colorscale=[[0.0, T.HIGHLIGHT_2], [0.5, "#F5F5F5"], [1.0, T.HIGHLIGHT]],
        zmid=100,
        hovertemplate="%{y}<br>%{x}: %{z:.1f}<extra></extra>",
        colorbar=dict(title="Index", thickness=12),
    ))
    T.titled(
        fig,
        "Global hubs and domestic ports keep almost the same weekly rhythm",
        "Mean daily port calls by weekday, indexed to each class's own weekly av",
        "Source: IMF PortWatch",
    )
    fig.update_layout(height=380)
    fig.show()
else:
    print("No vessel_count_total column in the reference table - inspect ports_r
    print(list(ports_ref.columns))
```

Classification column: port_class

Global hubs and domestic ports keep almost the same weekly rhythm

Mean daily port calls by weekday, indexed to each class's own weekly average = 100



What I found. This is a null result, and worth stating as one rather than forcing a story onto it. Globally systemic ports (the top 20 by vessel traffic) swing across the week by about **13 index points** peak-to-trough; regional/domestic ports swing by about **14.5 points** — both dip on weekends and both peak early in the week, and the gap between them is small enough that it could plausibly be noise rather than a genuine operational-tempo difference. The hypothesis that global hubs run closer to 24/7 while domestic ports keep a weekly rhythm is not supported by port-call counts at this resolution; it may still be true of *berth utilisation* or *dwell time*, neither of which this dataset captures — port calls measure vessel arrivals, not how continuously the port operates between them.

Q10 · Does the Fed's supply-chain index track the ships?

Why this question. The New York Fed's GSCPI is a widely watched composite built from PMI survey components and freight rates. Our data is physical. Do they agree, and does one lead the other? A survey-based index that lags the observable movement of ships by several weeks would be a genuinely useful finding.

Before running: download `gscpi_data.xlsx` from the [NY Fed](#) into `data/raw/`.

Method note. GSCPI is monthly, our data is daily — resample to month-end before comparing, and standardise both so they share an axis.

```
In [21]: GSCPI_PATH = D.RAW / "gscpi_data.xlsx"

if not GSCPI_PATH.exists():
    print("gscpi_data.xlsx not found - download from")
    print("https://www.newyorkfed.org/research/policy/gscpi")
```

```

else:
    g = pd.read_excel(GSCPI_PATH, sheet_name=0)
    print("Sheet columns:", list(g.columns)[:8])
    # The sheet has a few header rows - adjust the slice after inspecting it
    g = g.rename(columns={g.columns[0]: "date", g.columns[1]: "gscpi"})
    g["date"] = pd.to_datetime(g["date"], errors="coerce")
    g = g.dropna(subset=["date", "gscpi"])[["date", "gscpi"]]

    phys = (chokepoints.groupby("date", as_index=False)["n_total"].sum()
            .set_index("date")["n_total"].resample("ME").mean())
    comp = pd.DataFrame({"physical": phys}).join(
        g.set_index("date")["gscpi"].resample("ME").mean(), how="inner").dropna()

    z = (comp - comp.mean()) / comp.std()

    fig = go.Figure()
    fig.add_trace(go.Scatter(x=z.index, y=z["gscpi"], name="GSCPI (survey-based)",
                             line=dict(color=T.CONTEXT, width=2)))
    fig.add_trace(go.Scatter(x=z.index, y=z["physical"], name="Chokepoint transi",
                             line=dict(color=T.HIGHLIGHT, width=2.6)))
    T.titled(
        fig,
        "The survey-based pressure index and the observed fleet tell related but",
        "Both series standardised (z-scores), monthly",
        "Sources: IMF PortWatch; Federal Reserve Bank of New York",
    )
    fig.update_layout(height=460, yaxis_title="Standard deviations", xaxis_title="Date")
    fig.show()


    # Lead-lag: which shifts of GSCPI best correlate with the physical series?
    lags = range(-6, 7)
    corr = [{"lag_months": L,
             "corr": z["physical"].corr(z["gscpi"].shift(L))} for L in lags]
    corr = pd.DataFrame(corr).dropna()
    best = corr.loc[corr["corr"].abs().idxmax()]
    print(f"\nStrongest correlation at lag {int(best['lag_months'])} months: r =

```

Sheet columns: ['Date', 'GSCPI']

The survey-based pressure index and the observed fleet tell related stories

Both series standardised (z-scores), monthly



Strongest correlation at lag 6 months: $r = 0.25$

What I found. The best-fitting lag is **+6 months, and only weakly correlated ($r = 0.25$)**

— today's chokepoint transit volume correlates best with the GSCPI value from six months earlier, not the other way around. That is the opposite of what the "ships know before the surveys do" hypothesis predicted, and the correlation is weak enough ($r = 0.25$, on monthly data with limited independent observations) that it should not be over-read: it is plausible this reflects both series sharing a slow-moving macro trend (the post-COVID normalisation) rather than a genuine causal lead-lag relationship. The two lines in the chart above visibly diverge around the 2023-24 crisis — GSCPI, built from PMI survey components and freight rates, barely reacts to the chokepoint-level disruption that the AIS data shows clearly. The practical takeaway is closer to *the survey-based index and the physically observed fleet are measuring related but distinct things*, not that one reliably anticipates the other.

Conclusions

The shock was real and large. Suez and Bab el-Mandeb transits fell 62-76% from baseline and were still running at roughly half their pre-crisis level as of the most recent data — this was not a brief dip.

It was a rerouting, not a contraction, and it hit cargo types unevenly. The Cape of Good Hope did not just absorb the lost traffic, it more than doubled its own baseline volume (Q1). But the reroute was not uniform across cargo classes: Ro-Ro traffic

collapsed hardest (-93%), general cargo barely moved (-40%) — contrary to the initial hypothesis that container traffic would flee first (Q2).

It was unusually hard to escape. For 6.5 straight months, both of the Suez alternative's supporting routes were constrained at once — Panama Canal transits were down 53% from the drought at the same time Suez/Bab el-Mandeb were down 62-76% (Q4).

The costs landed unevenly and along a clear geographic line. Red Sea-facing ports (King Abdullah, Aqaba, Jeddah) lost 37-81% of container traffic while Mediterranean and Black Sea ports gained (Q5); global container concentration itself stayed flat at ~30-31% in the top 20 ports throughout (Q6) — the network's structural exposure to those twenty hubs did not change, even though the crisis passed through it.

Resilience is not evenly distributed, but the evidence for *why* is thinner than a tidy story would suggest. LPI and recovery speed correlate only weakly ($r = 0.39$, $n = 17$, with Singapore itself the biggest outlier) (Q7); the weekday-rhythm difference between global hubs and domestic ports that would support an "operational tempo" story turned out to be a null result (Q9); and import-dependence cleanly separates small-island economies from resource exporters but is a trade-balance proxy, not a measured routing exposure (Q8).

The physical record and the survey-based index tell related but distinct stories. GSCPI correlates only weakly with the observed chokepoint data ($r = 0.25$, best lag +6 months) — not strongly enough to claim either one reliably leads the other (Q10). Taken together: the disruption was real, it moved trade rather than destroying it, it compounded across two routes simultaneously, and it left clearly geographic winners and losers — but several of the "resilience" questions this analysis hoped would yield clean answers (LPI, port rhythm, the survey index) instead produced honest, modest, or null results, which is itself a finding about how hard resilience is to measure from AIS data alone.

Limitations

- PortWatch trade volumes are **AIS-derived estimates**, not customs data. All findings are framed as relative change for this reason.
 - Vessel *counts* are not *capacity*; a shift toward larger ships would understate volume relative to calls.
 - Event dates marking the crises come from public reporting and are context, not data.
 - Q7's relationship is correlational across a modest sample with a single shock — it cannot establish causation.
 - The chokepoint exposure proxy in Q8 is a documented modelling choice, not a measured routing assignment.
-

Sources

- **IMF PortWatch** — portwatch.imf.org · International Monetary Fund
- **World Bank Logistics Performance Index** — lpi.worldbank.org
- **NY Fed Global Supply Chain Pressure Index** — newyorkfed.org/research/policy/gscpi
- **UNCTAD Port Liner Shipping Connectivity Index** — unctadstat.unctad.org